TONGZHOU MU

Phone: (+1) 858-257-7603 | Email: tongzhou.mu@gmail.com | Web: http://cseweb.ucsd.edu/~t3mu/

EDUCATION

University of California, San Diego
- *Ph.D. in Computer Science & Engineering* — ***Advisor:*** *Hao Su* — 2019 – Present
- *M.S. in Computer Science & Engineering* — ***GPA:*** *4.0/4.0* — 2017 – 2019

Zhejiang University *B.Eng. in Computer Science* — ***GPA:*** *3.8/4.0,* ***Major GPA:*** *3.97/4.00* — 2013 – 2017

RESEARCH INTERESTS

- Reinforcement Learning, Concept Discovery and Reasoning, Robot Learning and Embodied AI

PUBLICATIONS

- **Tongzhou Mu***, Zhan Ling*, Fanbo Xiang*, Derek Yang*, Xuanlin Li*, Stone Tao, Zhiao Huang, Zhiwei Jia, and Hao Su. "ManiSkill: Generalizable Manipulation Skill Benchmark with Large-Scale Demonstrations." In *the 35th Conference on Neural Information Processing Systems (NeurIPS) Datasets and Benchmarks Track*, 2021.
- **Tongzhou Mu***, Jiayuan Gu*, Zhiwei Jia, Hao Tang, and Hao Su. "Refactoring Policy for Compositional Generalizability using Self-Supervised Object Proposals." In *the 34th Conference on Neural Information Processing Systems (NeurIPS)*, 2020.
- Stone Tao*, Xiaochen Li*, **Tongzhou Mu***, Zhiao Huang, Yuzhe Qin, and Hao Su. "Abstract-to-Executable Trajectory Translation for One-Shot Task Generalization." Submitted to *the 11th International Conference on Learning Representations (ICLR)*, 2023.
- **Tongzhou Mu**, Kaixiang Lin, Feiyang Niu, Govind Thattai. "Learning Two-Step Hybrid Policy for Graph-Based Interpretable Reinforcement Learning" Submitted to *Transactions on Machine Learning Research (TMLR)*, 2022
- Fangchen Liu, Zhan Ling, **Tongzhou Mu**, and Hao Su. "State Alignment-based Imitation Learning." In *the 8th International Conference on Learning Representations (ICLR)*, 2020.
- Xingchao Liu*, **Tongzhou Mu***, and Hao Su. "Transfer Value or Policy? A Value-centric Framework Towards Transferrable Continuous Reinforcement Learning." In *the Deep Reinforcement Learning Workshop at the 32th Conference on Neural Information Processing Systems (NeurIPS)*, 2018.
- Zebang Shen, Hui Qian, **Tongzhou Mu**, and Chao Zhang. "Accelerated Doubly Stochastic Gradient Algorithm for Large-scale Empirical Risk Minimization." In *the 26th International Joint Conference on Artificial Intelligence (IJCAI)*, 2017.
- Zebang Shen, Hui Qian, Tengfei Zhou, and **Tongzhou Mu**. "Adaptive Variance Reducing for Stochastic Gradient Descent." In *the 25th International Joint Conference on Artificial Intelligence (IJCAI)*, 2016.

** indicates equal contribution*

INDUSTRY EXPERIENCES

Amazon Alexa AI, *Full-Time Applied Scientist Intern* — Sunnyvale, United States
- Project: Interpretable RL for Text-Based Games with Graph Inputs — June 2021 – Sep 2021

Wormpex AI Research, *Full-Time Research Intern* — Seattle, United States
- Project: Store Layout Optimization based on Customer Behavior Model — June 2020 – Sep 2020

Intel AI, *Full-Time Research Intern* — San Diego, United States
- Project: Memory-Constrained Navigation via Combining RL and Planning — July 2019 – Sep 2019

Microsoft Research Asia, *Full-Time Intern at Visual Computing Group* — Beijing, China
- Project: Indoor Visual Navigation by Deep RL — Apr 2017 – Aug 2017

AWARDS & HONORS

- ACM-ICPC (International Collegiate Programming Contest) Asia Regional Contest **Gold Medal** — 2015
- China Computer Federation Elite Collegiate Award (top 108 in China) — 2016

PROFESSIONAL SERVICES

- Leading Organizer of *ICLR 2022 Workshop "Generalizable Policy Learning in the Physical World"*
- Conference and Journal Reviewer: *ICLR, ICCV, ICRA, IROS, AAAI, NeurIPS, RA-L*